

AI • SME

GUIDE & STRATEGY PAPER • EXECUTIVE LEADERSHIP

Start where it counts

The complete guide to how a mid-sized company adopts AI without spreading itself thin: **start small** where it hurts most, prove the value — and only then scale out. From an honest assessment of where you stand through team, use cases, pilot, law, change, tools and numbers to day-to-day operation — with concrete examples and ready-to-use templates.

AUDIENCE

Executive leadership • Project lead • Departments

SUBJECT

Adopting AI in mid-sized companies

AS OF

June 2026

ABOUT THE AUTHOR

Maris Reinold

Master of Engineering (M. Eng.) · 17 years of professional experience

Maris Reinold is an engineer (M. Eng.) with 17 years of professional experience — gained across several companies and in a range of departments and roles. This guide grew out of that broad view of technology, processes and people: a clear, workable path to adopting artificial intelligence in mid-sized companies with a sense of proportion — start small, prove the value, and consider technology, law and people together from day one.

| *“Start where it counts.”*

Contents

1	What this is about	4
2	Principles	6
3	How we proceed	7
4	An honest start: where we stand	9
5	Who drives it	11
6	Where it pays off	13
7	The first step: the pilot	16
8	Scaling out	18
9	Staying with it	19
10	What protects us: law and risk	20
11	Bringing people along	23
12	Tools	24
13	What it costs, what it returns	26
14	Where it usually fails	28
15	The first 90 days	29
16	The decision	31
17	Appendix — Templates & Checklists	32
A	Use-case profile	32
B	Core team & responsibilities (RACI)	32
C	Pilot checklist	33
D	Legal & compliance checklist	33
E	Kick-off workshop (agenda)	33
F	Tool evaluation scorecard	34
G	ROI worksheet	34
H	Cheat sheet: good AI prompts	34
I	Building blocks of a works agreement	35
J	Glossary	35

CHAPTER 1

What this is about

There is a lot of talk about AI right now and very little arithmetic. In most organisations a licence sits somewhere, a few colleagues are trying things out, and by the end of the quarter no one can say exactly what came of it. That is not a technology problem. It is the result of starting a little bit everywhere and nowhere properly.

This guide proposes something different, and it is deliberately immodest in its simplicity. We do not roll AI out everywhere at once, but in a single place — the one where it pinches most today. There we prove within a quarter whether it holds up. Only then do we scale out. That sounds cautious, but it is the opposite: it is the fastest way to a solid result, instead of experimenting a little on many fronts for a year.

The through-line in one sentence: a small, mandated team gathers the real pain points in the organisation, starts a tightly scoped pilot where the pressure is greatest, measures success against a number set in advance, and then scales out in waves — accompanied by training, clear ground rules and a sound legal framework.

THE KEY MESSAGES AT A GLANCE

- We start in **exactly one place** — where the pressure is greatest — rather than a little everywhere.
- A **small team with a clear mandate** drives it, not a many-headed committee.
- A **90-day pilot** proves the value against a number set in advance.
- Only after the proof do we **scale out in waves**, with training and clear rules.
- **Law, data protection and people** are considered from day one, not at the end.
- The project becomes a **lasting capability**, not a flash in the pan.

The difference between an expensive experiment and a good investment is where you begin.

For the executive team this comes down to three manageable commitments. One named person with a mandate who owns the topic. A small budget for a proper tool and training. And ninety days for a pilot, for which we set a number in advance to be measured against. At the end stands a simple decision — expand or stop — made on results, not on the mood in the building.

What is deliberately missing matters as much as what is here. No company-wide rollout, no collection of tools for every department, no year of lead time. All of that can come, but only once the first

step has proven itself. Begin broad and you spread your money across many small attempts, none of which grows large enough to convince. Begin narrow and after a quarter you have either a proof or a cheap lesson.

This guide is ordered the way such an undertaking unfolds, and it goes into detail. A few principles and an overview are followed by the individual stages — from the honest assessment of where you stand through team, use cases and pilot to scaling out and day-to-day operation. Then the cross-cutting themes that decide success or failure: law and risk, people, tools and money. A chapter on the most common mistakes, a concrete ninety-day plan and an appendix of ready-to-use templates and checklists round it off.

CHAPTER 2

Principles

Before we get to stages, tools and roles, a shared framework helps. The following principles are the compass: when a decision is hard, they point the way. None of them is new or original, but together they explain the difference between initiatives that move the needle and those that fizzle out.

Value before technology. We start with a real problem, not a tool. The question is always which work becomes easier, faster or better — not everything the tool can do.

Start small, learn fast. Short learning loops beat grand plans. A scoped pilot yields insights in weeks that outdo any months-long upfront analysis.

Start where it hurts. The first use belongs in the area with the highest pressure that AI can actually relieve. That is where willingness to change is greatest and success most visible.

A mandate is mandatory. AI changes ways of working and crosses departmental lines. Without visible backing from leadership — with budget and time — every initiative stays piecemeal.

Enable, don't mandate. People use what they understand and can operate safely. Training and support are not an add-on but the highest-return lever.

Work with reality. AI is no substitute for order. Where processes are unclear and data is messy, you start pragmatically and use AI to help create order, rather than waiting for the perfect state.

Law from the start. Data protection, confidentiality, co-determination and the requirements of AI law are considered from the beginning. That is not a brake but the precondition for lasting operation.

Make it measurable. Before every pilot comes the baseline number. What isn't measured can be neither steered nor defended, and a success without proof will be doubted.

People stay responsible. AI delivers drafts and suggestions; the professional and legal responsibility rests with people — especially where there is an external effect or personal data.

Not a project, but a capability. The undertaking has a beginning but no end. Models, tools and rules change quickly, so we build something lasting rather than a flash in the pan.

CHAPTER 3

How we proceed

The path can be captured in a few stages. They don't run strictly one after another; the first two overlap, and two themes run through everything — namely law and dealing with people. But as a map, the sequence works well.



Adoption proceeds in stages — from taking stock through the pilot to scaling out and day-to-day operation. Each stage builds on the previous one and leads a little further upward.

At the start stands a sober inventory: where do we stand on data, technology and above all people, and what are staff already using on their own. Almost in parallel we set up a small, capable team and give it a clear mandate. Then we gather not “AI ideas” but pain points, and sort them by value, effort and risk. From this list comes a single pilot, where the pressure is greatest. If it proves out, we scale out in waves, with training and a lean approval process for new applications. And finally the project becomes a permanent state — a capability we maintain.

Stage	Duration	Central question	Outcome
Taking stock	2–4 weeks	Where do we stand on data, technology, skills and rules?	Honest maturity picture, target vision, first in-house rule.
Team & mandate	in parallel	Who drives it, with what authority?	A small line-up with names and backing.
Gather pain points	2–3 weeks	Where is the greatest value at acceptable risk?	A sorted list, one or two pilot candidates.
Pilot	6–12 weeks	Does it work, and can we prove it?	Proof in numbers, a decision at the end.
Scaling out	3–9 months	How do we bring the value safely into the wider organisation?	Use across several areas, with ground rules.
Staying with it	ongoing	How do we stay effective, safe and current?	AI as a maintained capability rather than a one-off.

Two things run through all the stages because they decide success and failure. One is governance — clear rules, an overview of the applications in use and a lean approval path; the other is dealing with people, from explaining the why to visibly celebrating the first wins. Both get their own chapter later but must be considered from the first week.

Between the pilot and scaling out lies a deliberate stop where a decision is made: expand, sharpen or stop. This single point spares the most common misery of such undertakings — the pilot that never ends and ties up money without ever proving anything.

CHAPTER 4

An honest start: where we stand

Before you set off, you should know from where. Taking stock need not be a month-long project; half a day with leadership, IT and two or three department heads is enough for a solid first picture. The honest gaps will show up in the pilot at the latest anyway. It is not about a high score but about clarity on where we stand and where things stick.

Maturity across five dimensions

Five fields deserve a sober look. You can rate them roughly, from “barely present” to “established”, and quickly see where the biggest gaps lie.

Dimension	What it is about	Typical gap in mid-sized companies
Strategy & leadership	Is there a goal, sponsorship and budget? Is AI tied to the business?	AI is seen as a tech playground, not a business topic; no one is clearly responsible.
Data	Where does the relevant knowledge sit? Is it findable, current, accessible?	Knowledge sits in people’s heads, inboxes and folders; no searchable base.
Technology & IT	User management, security level, how well systems connect.	Lean IT with little slack; access is inconsistent; integrations are costly.
Skills & culture	Basic understanding, openness, learning culture, dealing with change.	A wide range: a few enthusiasts alongside widespread scepticism, little practice.
Rules & law	Are there usage rules? How mature is data protection? Is the works council in the loop?	No in-house rule; unofficial use via private accounts; works council left out.

Taking stock

Three things are looked at concretely. The data landscape — which sources exist, from the file store through the CRM to technical documentation, and how quality and access stand; this later decides which applications are realistic at all. The technical and security base, from central sign-in through the permissions concept to the question of what counts as confidential here. And the practice already being lived.

This last point deserves special attention because it is almost always underestimated. In virtually every company, staff are already using AI, often via private accounts and without IT’s knowledge. That is no scandal but a signal: there is real demand here. This unofficial use yields the most honest use cases of all. At the same time it carries a real risk, because anyone who enters confidential customer or personnel data into a free service can breach data protection and trade-secret law and let content flow uncontrolled into external models. The most important immediate measure is

therefore almost banal: a short, provisional rule on what is allowed, what is not and which approved tool we use.

What we take from this stage fits on one page: a picture of our maturity with the biggest gaps, a concise target vision for the next six to twelve months, a rough frame for time and money, and that first rule. No more is needed. Whoever waits for the perfect data state waits forever; it is enough to begin well enough and get better along the way.

CHAPTER 5

Who drives it

This is not a job for a committee. A twelve-person steering group that meets monthly and decides nothing is the surest way to let the whole thing fall asleep. What works is a small line-up with clear tasks and a real mandate that aligns briefly and often and is allowed to decide. Up to around 250 employees, three or four core roles are often enough, supplemented by rotating contacts from whichever department is affected.

The roles at the core

One responsible person drives the topic, steers the pilot and reports to leadership; that is the most important appointment, and it needs a name, not a job description. Alongside comes someone from the affected department who actually knows the work and uses the tool daily, because only then does it take root in everyday work. Someone from IT owns access, integration and the unobtrusive security questions. And data protection belongs in early, not only at the end when everything is built and changes are expensive. Where a works council exists, it too belongs in the information flow from the start. Larger organisations add a person for communication and training — the larger the company, the more important.

Who does what

So that responsibility doesn't stay vague, a simple mapping helps. *R* stands for the responsibility to carry it out, *A* for the decision and accountability (always exactly one place), *C* for “consulted” and *I* for “informed”.

Task	Exec/Sponsor	AI lead	IT	Data prot.	Dept.	WC
Strategy & budget	A	R	C	C	I	I
Pain points & prioritisation	I	A/R	C	C	R	I
Tool selection & security	I	C	A/R	C	C	I
Data protection, contract, impact assessment	I	C	C	A/R	I	C
In-house rule / works agreement	A	R	C	R	I	C
Run & measure the pilot	I	A	C	I	R	I
Training & communication	I	A	I	C	R	I
Decision to expand	A	R	C	C	C	I

A clear, visible mandate from the top is the cheapest accelerator there is — and its absence the most common reason for standstill.

The rhythm is simple: the core team aligns briefly each week, decision-oriented and in half an hour; once a month a larger group with leadership looks at priorities, budget and the upcoming decisions. Leadership's task in this setup is clearly defined and small: decide, release the budget, back the responsible person, and talk about the result after ninety days. What's needed is visible backing, best captured in two or three written sentences. Without such a mandate the initiative becomes everyone's side project and dies of jurisdictional questions.

Once the pilot proves out and the thing scales, a network of contacts in the departments comes later — open-minded colleagues who help first and act as role models. That is the most important support for broad acceptance, but a matter for the second stage, not the start. At the beginning: as few people as possible, as much mandate as necessary.

CHAPTER 6

Where it pays off

Now it gets concrete. The best question to put to the departments is not “Where could we use AI” but “Which work here is repetitive, tedious, time-consuming or error-prone”. You ask about pain points, not technology; translating them into a solution is the team’s job. A short workshop per area, a few conversations with experienced people, the occasional half-day shadowing daily work, and a low-threshold channel for suggestions bring together more in a few weeks than you can process. And the practice already lived unofficially is the most honest indicator of real demand.

Where AI typically moves the needle

Wherever work consists of language, research and data, it is strong. Roughly, it can be outlined like this, without slipping into enthusiasm.

Area	Typical application	Risk / note
Text & knowledge	Drafts, summaries, translations, minutes, research, search of your own knowledge	low — a good entry point
Customer service	Triage requests, draft replies, knowledge base, sentiment	low, if a person checks
Sales	Pre-contact research, draft quotes and outreach, CRM upkeep	low to medium; fact-check
Back office & finance	Read documents, reconcile, draft reports, search contracts	medium — four-eyes principle
Engineering & production	Documentation, research, drafting test plans, coding help	depends on data quality
IT & development	Coding and review help, scripts, documentation	high productivity lever
HR	Job ads, templates, research	Caution — high for selection/evaluation

The last row is a warning. AI that helps decide on hiring, evaluation, promotion or dismissal is legally sensitive and subject to strict requirements. Such cases are not suitable for getting started; you flag them and handle them separately later. Chapter 10 goes into this.

Concrete examples from the departments

So the list doesn’t stay abstract, here are a few concrete, well-measurable applications for the most important areas — as inspiration, not as a prescription.

Sales. Before each meeting, a two-minute customer brief from public sources and the CRM; draft quotes and outreach based on proven text blocks; transfer meeting notes into the CRM automatically. Effect: more time with the customer, less at the desk.

Customer service. Automatically triage incoming requests and add a draft reply from your own knowledge base, which a person checks and sends; condense recurring questions into a maintained knowledge base. Effect: shorter handling time, more consistent quality.

Back office & finance. Read invoices and receipts and pre-code them; reconcile figures from several sources; prepare draft reports; search long contracts specifically for deadlines and clauses. Effect: less typing, fewer transfer errors — with the four-eyes principle.

Engineering & production. Draft technical documentation and test plans; quickly find the right passage in standards and manuals; generate recurring analyses and small scripts. Effect: knowledge becomes findable and usable faster.

IT & development. Coding and review help, tests and documentation, preparing log data. Effect: one of the highest productivity levers there is, often double-digit time savings.

Marketing & communication. First text drafts for web, newsletters and social media; translations; image and topic ideas; condensing customer feedback. Effect: faster from idea to draft, more output with the same team.

Describe and prioritise

Every serious idea is described briefly and uniformly, ideally on half a page: the problem, the solution idea, the expected value, the data needed including any personal data, the effort and the risk. This makes proposals comparable and forces clarity. A ready-made template is in the appendix (A).

Assessment runs along four questions, on a simple scale: How great is the value? How easy the implementation? How high the current pain? And how low the legal risk? Value and feasibility weigh most. A simple four-box view orders the result.

	low effort	high effort
high value	Quick wins ideal pilot candidates	Big undertakings plan, tackle later
low value	Side effects when convenient	Time sinks avoid for now

For the start we deliberately choose from the top-left box: high value, low effort. The ideal first case also has a high pain level and can be measured cleanly, and it is a lighthouse — visible enough that the success makes the rounds in the organisation. Because the best argument for the next step is

not a slide but an area that was noticeably helped. At the end of this stage stands a sorted list and one or two candidates for the pilot, with the favourite where the pressure is greatest and solvable by AI.

CHAPTER 7

The first step: the pilot

The pilot is the centrepiece, and it deliberately belongs where it hurts most. The temptation is great to start with the open-minded team that is already experimenting anyway. More effective is the area with the highest pain. Whoever suffers from overload, backlog or constant firefighting is open to real relief, and that is exactly where the difference between before and after is most clearly measurable. A relieved team becomes a convinced advocate, and that one story later carries the rollout. It convinces more than any glossy demonstration.

Three conditions

For the approach to hold, three things must be right, otherwise you take the next-best candidate. The pressure must be solvable by AI, because not every problem is an AI problem; sometimes there is simply a lack of staff, processes or leadership. Despite the load, the team needs a little room to take part, otherwise the pilot becomes an extra burden. And it needs a department head who actively stands behind it.

Choosing the pilot

From the candidates of the previous stage comes exactly one, at most two. These criteria have proven useful:

Criterion	Why it counts
Pain level	High pressure that AI can relieve drives motivation and visibility.
Feasibility	Achievable in six to twelve weeks, without major integration.
Measurability	Clear before-values exist or can be gathered quickly.
Willing sponsor	The department head actively supports it and provides time.
Manageable risk	Ideally no high-risk field, limited personal data.
Clear scope	A tight frame: one or two applications, a small group.

Measure first, then start

The most important and most often skipped step: we measure the starting state before anything begins. Without a before-value, no success can be proven later, and an unprovable success will be doubted. This includes a simple one-sentence hypothesis (“We believe that [application] improves [metric] for [group] by [target]”), the before-values gathered, a target number above which the pilot counts as successful, and a stop criterion — because a clear no is also a valuable result.

No baseline, no proof. Whoever doesn't make success measurable beforehand won't be able to defend it afterwards.

Scope, tool, data protection

Tight scope means: one or two applications, a small representative group of usually five to fifteen people, a fixed period of six to twelve weeks. For the tool we take an enterprise version with data-protection guarantees, not a private free account. And a pilot is already data processing: we sign a contract with the provider, we check whether a data-protection impact assessment is needed, we feed in no confidential or personal data without a basis, we involve the works council and we train the pilot group. The simplest is to cut the pilot so it does without sensitive data. The law chapter spells this out.

Run and decide

A short introduction explains what the tool can and cannot do, how to phrase good prompts and which rules apply; a handful of prepared templates lowers the hurdle considerably. A brief weekly exchange clears obstacles immediately, the numbers run along, and every insight flows into a growing playbook.

An example makes it tangible. Take a service team drowning in requests; the backlog grows, morale sinks. Today a reply takes twelve minutes on average, forty emails are left each day. In the pilot, the system drafts the replies from our own knowledge base, and a person checks and sends. The number we want to beat: thirty percent less handling time at the same quality, used by at least four-fifths of the group. If the pilot hits that number, the debate about value is over, because it is settled on our own data. If it doesn't, we have learned a lot for little money.

SECOND EXAMPLE · QUOTES IN SALES

Creating a quote today ties up ninety minutes on average, because text blocks, prices and earlier quotes have to be laboriously gathered. In the pilot, the tool produces a complete draft from keywords and CRM data, which sales checks and adjusts. The number we want to beat: half the time per quote at the same hit rate — and as a bonus, quotes reach the customer faster, which experience shows raises the probability of closing.

At the end stands a deliberate cut, not a creeping extension: expand, because the number holds; sharpen, because it is promising but not yet ready; or stop and take the next bottleneck. And a successful pilot is only complete once its story is told, with before/after numbers and the voices of those involved. This communication is no add-on, it is the fuel for scaling. You'll find the concrete ninety-day sequence in Chapter 15, a checklist in the appendix (C).

CHAPTER 8

Scaling out

The pilot has convinced; now it goes into the wider organisation. The biggest mistake here is the big bang: switching everything on for everyone at once, without support. That creates overwhelm, frustration and ultimately rejection. Smarter is to proceed in waves. First the pilot area in full and one or two nearby departments with similar needs, then the next by priority, each with its own training and applications adapted to its everyday work, and each wave evaluated before the next starts.

The technical side

With scale come a few concrete tasks, all aimed in one direction: grow safely and manageably.

Building block	What to watch
Licences	Roll out to need rather than blanket; watch usage, reclaim unused access.
Access & identity	Sign-in via central user management, rights by role, clean joining and leaving.
Integration	Connect step by step to your own content, CRM and stores, where the value is greatest.
Guardrails	Data classification, blocks for sensitive data, logging.

Enablement and everyday help

What was personal in the pilot needs a little structure at scale. Training that differentiates by role, because sales needs different examples than accounting. The network of contacts from Chapter 5 that carries enablement into the departments. A small collection of proven templates that everyone can reach. And a simple, multi-level path where help arises: first self-help, then the contact(s) in the area, then the core team for the tricky cases. A shared place to exchange — a channel or a regular format — keeps the knowledge alive.

Keeping order

So that scale doesn't become a free-for-all, a lean approval path for new applications is needed. Whoever wants to start something new describes it briefly, the team checks value, data and risk and approves or escalates the question one level up. Everything approved lands in a simple overview of all AI applications in use, which is part of the mandatory documentation anyway. And for tools: better to consolidate on one or two that fit well and are used than to run a dozen licences, half of which lie idle and are still billed every month.

CHAPTER 9

Staying with it

The rollout does not end the undertaking; it only changes the mode. The project becomes an ongoing task, and the core team shifts from project mode into operating mode. Models, tools and rules change quickly; whoever sets up once and then stops looking has, within a year, an unwatched free-for-all.

The ongoing tasks are manageable: gather, assess and implement new use cases so the pipeline never runs dry; keep an eye on usage, value, quality and incidents; assess new models and tools soberly rather than chasing every trend; keep training and refreshing; and watch the legal situation, which is especially in motion right now.

What changes	How we respond
New models and features	Assess in a structured way rather than adopting reflexively; introduce only what brings real value.
Changed legal situation	A fixed responsibility for monitoring; keep the in-house rule and overview current.
Growing usage	Steer cost and licences; let help and approval processes grow with it.
New risks	Learn from mistakes, sharpen guardrails, maintain a simple reporting path.

Once a year an honest review pays off: what runs productively and delivers value, what was discontinued, how maturity has developed, where we stand on cost, skills and rules. This rhythm is the cheapest insurance against initial enthusiasm turning into an expensive, unmanageable patchwork.

CHAPTER 10

What protects us: law and risk

This chapter is an orientation, not legal advice; the binding questions we clarify with our data-protection function and, where appropriate, legal counsel. The legal position is as of June 2026, and AI law in particular is in flux. The good news upfront: for the typical entry cases the framework is manageable, as long as you consider it from the start rather than at the end. AI touches several areas of law at once; we go through them in turn.

European AI law

EU AI law sorts applications by risk: the higher the risk, the stricter the obligations. Four tiers can be distinguished.

Tier	Meaning	For us
Prohibited	Unacceptable practices such as social scoring of people or certain manipulative systems.	Not an issue for ordinary applications, but worth knowing.
High-risk	AI in sensitive fields, e.g. HR decisions, credit-worthiness or critical infrastructure.	Relevant e.g. for AI in HR decisions — extensive obligations, handle separately.
Limited risk	Systems with a transparency duty, e.g. chatbots or AI content that must be clearly labelled.	A frequent case; mind the labelling.
Minimal risk	All other applications, e.g. text assistance, summaries, research.	The bulk of entry cases.

By far the largest part of what we have planned falls into the lower two tiers. The practically most important obligations are then transparency — making AI use and AI content recognisable — and AI literacy, which we come to next. The only thing you must actively check is whether a case slips into the high-risk area, above all with HR decisions.

THE DUTY THAT ALREADY APPLIES TODAY: AI LITERACY

One requirement affects virtually every company immediately: whoever uses AI must ensure that the people involved understand it sufficiently. This has applied since February 2025. Translated, that means training is not optional but required, and it should be recorded with proof of attendance. Fortunately this coincides with what we should be doing anyway per Chapter 11.

THE DEADLINES, SOBERLY CONSIDERED

A lot of half-knowledge circulates here, hence the sober look. Originally the strict high-risk obligations were to take effect from August 2026. With the so-called Digital Omnibus, a postponement was set in motion; in May 2026 the Council and Parliament reached a provisional agreement whose formal adoption is expected in mid-2026.

Obligation / milestone	Originally	After Digital Omnibus*
Prohibited practices	since Feb 2025	unchanged, in force
AI literacy duty	since Feb 2025	unchanged, in force
Obligations for general-purpose AI models	since Aug 2025	unchanged, in force
High-risk (standalone systems)	Aug 2026	Dec 2027
High-risk (embedded in products)	Aug 2027	Aug 2028
Transparency and labelling duties	Aug 2026	Dec 2026

* Provisionally agreed values (as of June 2026); binding only upon formal adoption and publication. The basic architecture of the law remains unchanged.

For us this means two things. On the high-risk topics there is time, which we use for sound setup, not for waiting. And data protection, which is unaffected, applies fully from day one.

Data protection

As soon as we touch personal data — and that happens quickly — data-protection law applies in full. The most important points at a glance.

Topic	What to watch
Legal basis & purpose	Every processing needs a basis and a clear purpose; no repurposing.
Data minimisation	Only as much data as necessary; where pseudonymisation suffices, no clear data.
Data processing agreement	A corresponding contract with every provider that processes data on our behalf.
Third-country transfer	For providers outside the EU, clarify the transfer basis beforehand.
Impact assessment	For high risk to individuals, a data-protection impact assessment before starting.
No training on our data	Configure the tool so that inputs do not serve the provider's training.

The simplest everyday rule we give everyone: no personal or confidential data into free, private services.

Co-determination, brand, copyright and trade secrets

If we have a works council, we involve it early. Tools that could be used to evaluate employees' performance or behaviour are subject to co-determination; the works council may also bring in experts for assessment. A supported works agreement on AI creates clarity for all sides and turns a potential brake into a partner.

The biggest everyday risk is in any case not the supervisory authority but a wrong or embarrassing statement in our name. Hence the simple, hard rule: nothing goes to a customer unchecked. On copyright, purely machine-generated texts generally enjoy no protection of their own, and third-party protected material does not belong in the machine unasked; for generated program code it is

worth looking at licensing questions. And our trade secrets — from design data through calculations to customer lists — do not belong in unsuitable tools, because we can lose their protection by doing so. Information security includes access control, a clear data classification, no credentials or secrets in inputs, vetted providers and a defined path for incidents.

What we document

A few documents form the practical backbone of sound governance. They need not be thick, but they must exist.

Document	Purpose
AI usage policy	Binding rules: permitted and prohibited use, approved tools, handling of data.
Overview of applications	Where AI is used for what, with which data, what risk, whose responsibility.
Impact assessment	For risk-laden processing; records risks and protective measures.
Provider contracts	Data processing agreements with all relevant providers.
Works agreement	Governs co-determination and protection of employees, where a works council exists.
Training records	Evidence of the fulfilled AI-literacy duty.

Whoever observes these few things from the start — a clear in-house rule, an approved and data-protection-compliant tool, trained people, a human before every output with an external effect, and a maintained overview — has the framework under control. A compact checklist is in the appendix (D).

CHAPTER 11

Bringing people along

Technology is rarely the problem. Success or failure is decided by whether the workforce comes along. That is not soft background music but the greatest lever of all, and it can be named without drifting into psychology.

Beneath the surface lie a few concrete worries. The first, often unspoken: will this cost me my job. No reassurance helps here, only the visible practice that AI takes over routine load and frees people for what they do better. The second worry is about surveillance; you meet it with openness and clear rules, not with covert performance monitoring. The third is the quiet “I can’t do this”, especially among long-serving, deserving colleagues; it disappears when the first steps succeed on real tasks from their own everyday work and no one feels exposed. And the fourth, entirely justified, is distrust of the results; you take it seriously by making clear that people remain responsible and the machine delivers drafts.

Three things are effective at the core. Explain the why, again and again, because people follow a purpose, not a tool. Lead by example rather than decree: when a manager says in a meeting that they just prepared this with the assistant, that works more strongly than any instruction. And make successes visible, with the numbers from the pilot and the voices of those involved. A learning culture in which questions and mistakes are allowed is not a luxury here but the condition under which people try anything new at all.

Training with a sense of proportion

Training is both a duty and a lever here: it covers the legally required AI literacy and incidentally decides actual use. It makes sense to tier it rather than serve everyone the same thing.

Tier	For whom	Content
Basic understanding	all employees	What AI is, what changes, which rules apply, where to get help.
Application	active user groups	Practical use, good prompts, safe use, examples from their own area.
Contacts	multipliers	A little depth, ability to help, gathering and passing on feedback.
Leadership	managers	Strategy, responsibility, role-model function, AI in leadership work.

CHAPTER 12

Tools

Tools are a means to an end, but the choice deserves thought. Roughly, there are four kinds.

Kind	What it is	Typical use
Enterprise assistant	A centrally managed chat assistant based on capable language models.	Broad entry: text, knowledge, research, drafts.
AI in existing software	Built-in features in office, CRM, ERP or industry software.	Value in the familiar tool, low hurdle.
Specialist tools	Solutions for narrowly defined tasks such as document extraction or translation.	Recurring specialist tasks.
In-house development	Tailor-made, often connected to your own knowledge.	Where standard isn't enough and the value is high.

Buy or build

For the start a simple order applies: first buy and achieve impact with standard, build selectively where the value clearly justifies in-house development.

Buy when ...	Build when ...
... a standard solution covers the need well.	... the case is specific and at the same time valuable.
... fast impact and little effort count.	... your own knowledge must be deeply integrated.
... internal development capacity is scarce.	... differentiation and independence matter.

Why an enterprise version, not a private account

For business use we take an enterprise version as a rule. The difference lies not in the model's ability but in the things a company needs: central administration, sign-in via existing accounts, traceable usage and the contractual assurance that our inputs do not serve training. In selecting, we watch a few criteria.

Criterion	Guiding question
Data protection & security	EU processing possible? Contract? No training on our data? Recognised certification?
Integration	Can it be coupled to our systems, sign-in and knowledge sources?
Usability	Is it usable without long onboarding?
Total cost	Licence, rollout, training, operation — not just the list price.
Independence	Can the model be switched without rebuilding everything?

The last point deserves emphasis. The supposedly best model changes every six months; whoever ties themselves firmly to one pays for it later. Flexibility is today a criterion in its own right.

CHAPTER 13

What it costs, what it returns

What isn't measured can be neither steered nor defended. Measurement not only secures the budget, it steers the initiative to where it really works and protects against flashes in the pan and self-deception. Four fields are worth a look.

Field	Example metrics	What it tells you
Usage	active users, frequency, number of productive applications	Is it being used at all?
Value	time saved, throughput, handling time, volume	Does it bring the hoped-for value?
Quality	error rate, rework, satisfaction	Does quality hold or rise?
Risk	incidents, rule breaches, share of unofficial use	Does it stay safe and compliant?

Let's talk about money, honestly. The running costs of getting started are small and easy to calculate: an enterprise tool per person per month, a one-off effort to connect to your own content and systems, time for training, and a fraction of a role for the responsible person. The biggest item in the first year is not the licence but attention — the time of capable people. Whoever underestimates that buys licences and wonders why nothing happens.

WORKED EXAMPLE · SERVICE TEAM

Eight people, each gaining about four hours a week from draft replies — 32 hours a week, at around 45 working weeks roughly 1,440 hours a year. At a fully-loaded hourly rate of €50, that corresponds to a gross value of around €72,000 a year.

Against that stand licence and operating costs of perhaps €6,000 and a one-off rollout including training of about €8,000 — making, in the first year, a **net value in the order of €58,000**. Only hard effects are counted, and only if the time gained is actually used. A blank worksheet for this is in the appendix (G).

On the other side stands the value: more throughput with the same team, shorter handling times, less rework, work that previously stayed undone and now gets done, new capabilities. Here the honest distinction is worth making between hard effects that can be proven directly and soft ones that are real but hard to quantify, such as relief and satisfaction. You may name both, but you should not artificially convert the soft ones into euros.

A warning belongs here so we don't fool ourselves: time saved becomes value only if it actually flows into something more valuable. If it seeps away, it was never there. So per use case we measure

exactly one simple number everyone understands, such as the hours a team wins back per week. Leading indicators like usage and training rate quickly show whether we're on track; lagging indicators like the value actually realised confirm the success over time. If the pilot proves out against its number, the case for expansion writes itself from our own data.

CHAPTER 14

Where it usually fails

You learn more from failed undertakings than from successful ones. The patterns repeat, and they are as avoidable as they are common. At the core, almost all share the same root: starting too broadly and too arbitrarily instead of narrowly and decisively. The following overview sets each mistake against its remedy.

Common mistake	Remedy
Starting with the tool instead of the problem	First the pain point, then the means.
The big bang: everything for everyone at once	Pilot, then waves, each supported and evaluated.
Treating AI as a pure IT project	Set it up as a change initiative, with departments and leadership.
No mandate, no one named	Visible backing from the top, one named responsible person.
No baseline, no goal	Measure beforehand, set a target number, report regularly.
Law and works council only at the end	Consider data protection, AI literacy and co-determination from day one.
Ignoring unofficial use	Make it visible, offer an approved alternative, clear rules.
The pilot that never ends	A fixed stage gate with a clear choice: expand, sharpen, stop.
Promising too much (“AI replaces department X”)	Speak of relief and name the limits openly.
Rolling out tools without enablement	Plan training, templates and support in firmly.
Sensitive data in unsuitable tools	Enterprise tool, data classification, clear rules, technical blocks.
A free-for-all and chasing every model	Consolidate on a few tools, assess the new soberly.

CHAPTER 15

The first 90 days

From intent into action. The following sequence brings the first stages into a workable rhythm. It is a suggestion, not dogma; the goal of the first ninety days is simply a running, measurable pilot and a sound framework around it.

Days 1–30 · Set up and understand

Work package	Concrete steps
Secure sponsorship	Win leadership over, capture mandate and budget frame in a few written sentences.
Form the team	Fill three or four core roles, set the working rhythm and decision rights.
Take stock	A half-day maturity look; capture data, technology and the practice already lived.
Law & co-determination	Involve data protection and, if present, the works council; communicate a provisional in-house rule.
Assess tools	Evaluate one or two enterprise tools; prepare the contract.

Days 31–60 · Focus and prepare

Work package	Concrete steps
Gather pain points	Short workshops in the departments; fill in use-case profiles.
Prioritise	Assess by value, effort, pain level and risk; settle the list and candidates.
Choose the pilot	Pick the area with the greatest solvable pressure; set scope, group, period.
Measure	Gather before-values; set hypothesis, target number and stop criterion.
Finalise the legal side	Sign the contract, check the impact assessment, ensure works-council involvement.
Enable	Provide the tool, train the group, create templates.

Days 61–90 · Trial and decide

Work package	Concrete steps
Run the pilot	Productive use in the group; brief weekly alignment; clear obstacles immediately.
Track	Gather feedback, follow the numbers, keep the playbook going.
Evaluate	Hold the result against the target number; produce a before/after comparison.
Decide	A clear choice: expand, sharpen or stop.
Tell & plan	Share the success story in the organisation; prepare the wave rollout and the final in-house rule.

The value of this tight cadence lies in the fixed stage gate after ninety days. It forces a decision and prevents an experiment from becoming a permanent state that ties up money and proves nothing.

CHAPTER 16

The decision

We have shown how AI can be adopted in mid-sized companies without spreading yourself thin, who drives it, where it pays off, what to watch legally, how to bring people along and what it costs and returns. What is missing now is a decision. It is deliberately kept small, because the point of this approach is to gain clarity quickly with little outlay before larger sums are at stake.

WHAT WE ASK OF LEADERSHIP

- One **responsible person with a mandate** who drives the topic and reports.
- A **manageable budget** for an enterprise tool including training and the one-off integration.
- **90 days** for a pilot where the pressure is greatest, with a number set in advance as the yardstick.
- A **fixed date afterwards** at which we decide on the results: expand, sharpen or stop.

The next sensible step is not another tool and not another presentation. It is a pilot with a number to be measured against. If the result is good, it carries itself into the wider organisation. If it is poor, we have avoided an expensive illusion for little money. Both outcomes are better than the state many are stuck in: a lot of talk about AI and, at the end of the quarter, no one who can say what it delivered.

APPENDIX

Templates & Checklists

Ready-to-use tools to copy and adapt. These templates accompany the first stages and can be adjusted to your own circumstances.

A — Use-case profile

Title & area	Short name and affected department
Problem today	Which activity, how often, how time-consuming, how error-prone?
Solution idea with AI	What exactly should AI do? Who uses it how?
Expected value	Time / quality / throughput / satisfaction — with a number where possible
Data needed	Sources · personal data (yes/no) · confidentiality
Effort / feasibility	Tool available? Integration needed? Complexity (low/medium/high)
Risk / law	High-risk? Data protection? External effect? Co-determination relevant?
Assessment (1–5 each)	Value __ · Feasibility __ · Pain level __ · Risk (low = 5) __
Decision	☐ Pilot candidate ☐ List ☐ Deferred ☐ Rejected

B — Core team & responsibilities

- ☐ **Responsible person (AI lead)** — Name: _____ · Mandate & time share clarified?
- ☐ **Department contact** — Name: _____ · per use case
- ☐ **IT & security** — Name: _____
- ☐ **Data protection** — Name: _____
- ☐ **Works council contact** — Name: _____ · Information flow in place?
- ☐ **Communication/training** (optional, larger organisations) — Name: _____
- ☐ **Rhythm** — weekly alignment · monthly with leadership
- ☐ **Decision rights** — clarified what the team decides on its own

The mapping of “who does what” (R/A/C/I) is shown as a table in Chapter 5.

C — Pilot checklist

- ☐ **Area chosen** — greatest solvable pressure · willing sponsor · some capacity
- ☐ **Scope defined** — 1-2 applications · 5-15 people · 6-12 weeks
- ☐ **Hypothesis & target number** — set and confirmed by the sponsor
- ☐ **Before-values gathered** — before-numbers documented
- ☐ **Tool ready** — enterprise version · sign-in · no training on our data
- ☐ **Legal done** — contract · impact assessment checked · works council involved · no sensitive data
- ☐ **Enablement** — introduction · training (proof) · templates · office hours
- ☐ **Support** — weekly alignment · feedback & numbers · playbook
- ☐ **Stage gate** — evaluation · clear choice: expand / sharpen / stop
- ☐ **Communication** — success story prepared and shared

D — Legal & compliance checklist

- ☐ **In-house rule** on AI use created and communicated
- ☐ **Overview** of applications in use created and maintained
- ☐ **Risk tier** classified per application (prohibited / high / limited / minimal)
- ☐ **AI literacy** — training organised, records documented
- ☐ **Transparency** — AI use and AI content made recognisable
- ☐ **Provider contract** (data processing agreement) signed
- ☐ **Third-country transfer** clarified
- ☐ **Impact assessment** — necessity checked, carried out if needed
- ☐ **Co-determination** — works council involved, works agreement if applicable
- ☐ **Security & secrets** — data classification, access, reporting path, no secret into the wrong tool

E — Kick-off workshop (agenda, approx. 90 minutes)

Duration	Topic	Content
10 min	Why & goal	The point of it, what the team gains, the picture of success.
15 min	Ground rules	Allowed and not allowed, data protection, people stay responsible.
30 min	Try it yourselves	Solve real tasks of the area together, practise good prompts.
15 min	Templates	Present and test area-specific templates.
10 min	Process & measurement	Period, alignment, what is measured, where to get help.
10 min	Questions & reassurance	Room for concerns; name contacts and office hours.

F — Tool evaluation scorecard

A rating from one to five per criterion, multiplied by the weight, gives the points. The sum makes tools comparable — but the conversation about it remains decisive.

Criterion	Weight	Rating 1-5	Points
Data protection & security	× 3	—	—
Integration with our systems	× 2	—	—
Usability	× 2	—	—
Total cost (TCO)	× 2	—	—
Independence / model switch	× 2	—	—
Feature scope for our cases	× 1	—	—
Total			—

G — ROI worksheet

Use case	_____
People affected	_____ people
Time saved per person & week	_____ hours
Total hours per week	People × hours = _____
Hours per year	× approx. 45 weeks = _____
Hourly rate (fully loaded)	_____ €
Gross value per year	Hours × rate = _____ €
minus licence & operation / year	- _____ €
minus rollout & training (year 1)	- _____ €
Net value year 1	= _____ €

Count only hard, provable effects — and only if the time gained is actually used.

H — Cheat sheet: good AI prompts

Give a role. Tell the AI who it should be — e.g. “You are an experienced sales assistant.” This shapes tone and depth.

Name the goal and format. What should come out, in what form (bullet points, email, table) and how long?

Provide context. Paste the relevant information directly — the customer’s text, key figures, an earlier example. AI can only work with what it sees.

Show an example. A good model works more strongly than any description.

Refine step by step. First a draft, then sharpen deliberately: “shorter”, “more formal”, “with reference to quote X”.

Always check. Verify facts, figures and names — responsibility stays with people.

WEAK → BETTER

Weak: “Write an email to a customer about a delay.”

Better: “You are a service representative. Write a polite, short email to Ms Müller: her delivery is delayed by three days due to a supply shortage. Apologise, give the new date (Friday) and offer direct contact for questions. Tone: assured, not servile.”

I — Building blocks of a works agreement

- **Preamble & purpose** — the aim of the agreement and a shared understanding.
- **Scope** — which tools, areas and groups of people.
- **Permitted and prohibited use** — with clear examples.
- **Exclusion of covert performance and behaviour monitoring.**
- **Data protection & data minimisation** — permissible data types, no training on our data.
- **Transparency & involvement** — information, disclosure, participation of the works council.
- **Qualification** — entitlement to training before use.
- **Mistakes & incidents** — reporting path, no sanction for honest mistakes.
- **Review & term** — regular evaluation, rules for adjustment and termination.

J — Glossary

Term	Short explanation
EU AI law (AI Act)	EU regulation that regulates AI by risk.
Digital Omnibus	EU package to simplify and postpone parts of AI law (in adoption, as of June 2026).
AI literacy	The duty to ensure users have a sufficient understanding.
Data processing agreement	A contract with a provider that processes data on our behalf.
Impact assessment	An examination of data-protection risks before risk-laden processing.
Baseline	The state measured before the start, for the before/after comparison.
Stage gate	A deliberate decision point: expand, sharpen or stop.
Unofficial use	AI use without approval, often via private accounts.

A practical orientation, not legal advice. As of June 2026.